

WHERE'S WALDO?

By : Deautaun Ross, Matilda Sorić, Nayma Kim

Outline



Project Goal



Data



Analyzing



Model



Conclusions

Project Goal

Find Waldo with
Image Recognition



Data



Images

- Waldo
- Not Waldo
- Original Images

From

- Found on Kaggle
[Kaggle Link](#)

File Size

376.93 MB



Analyzing



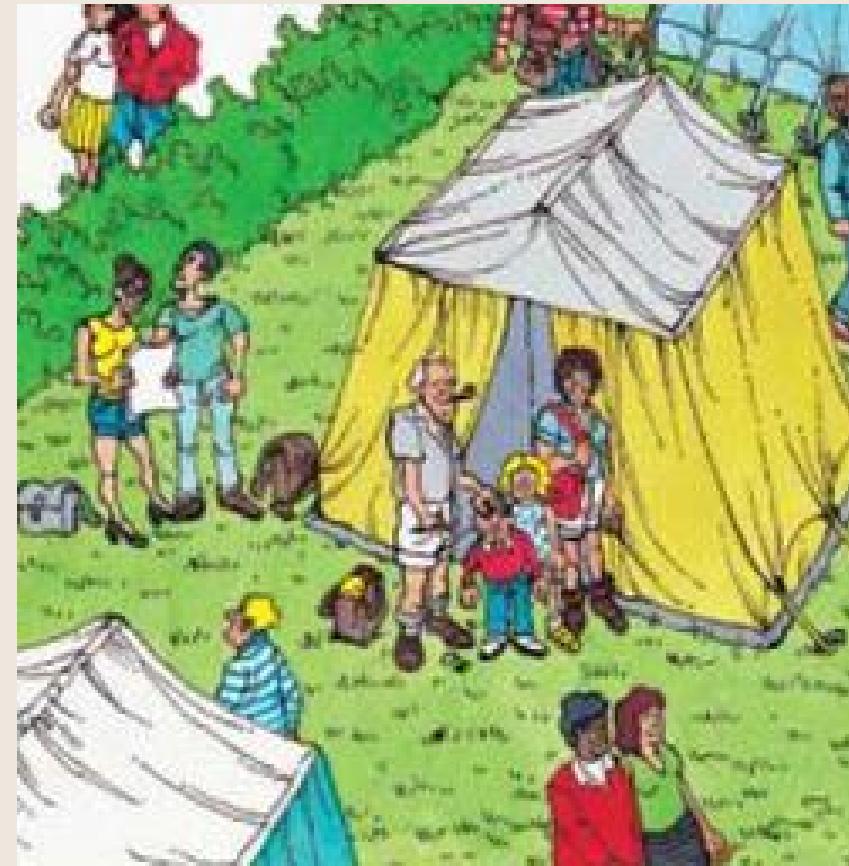
Waldo

39 images



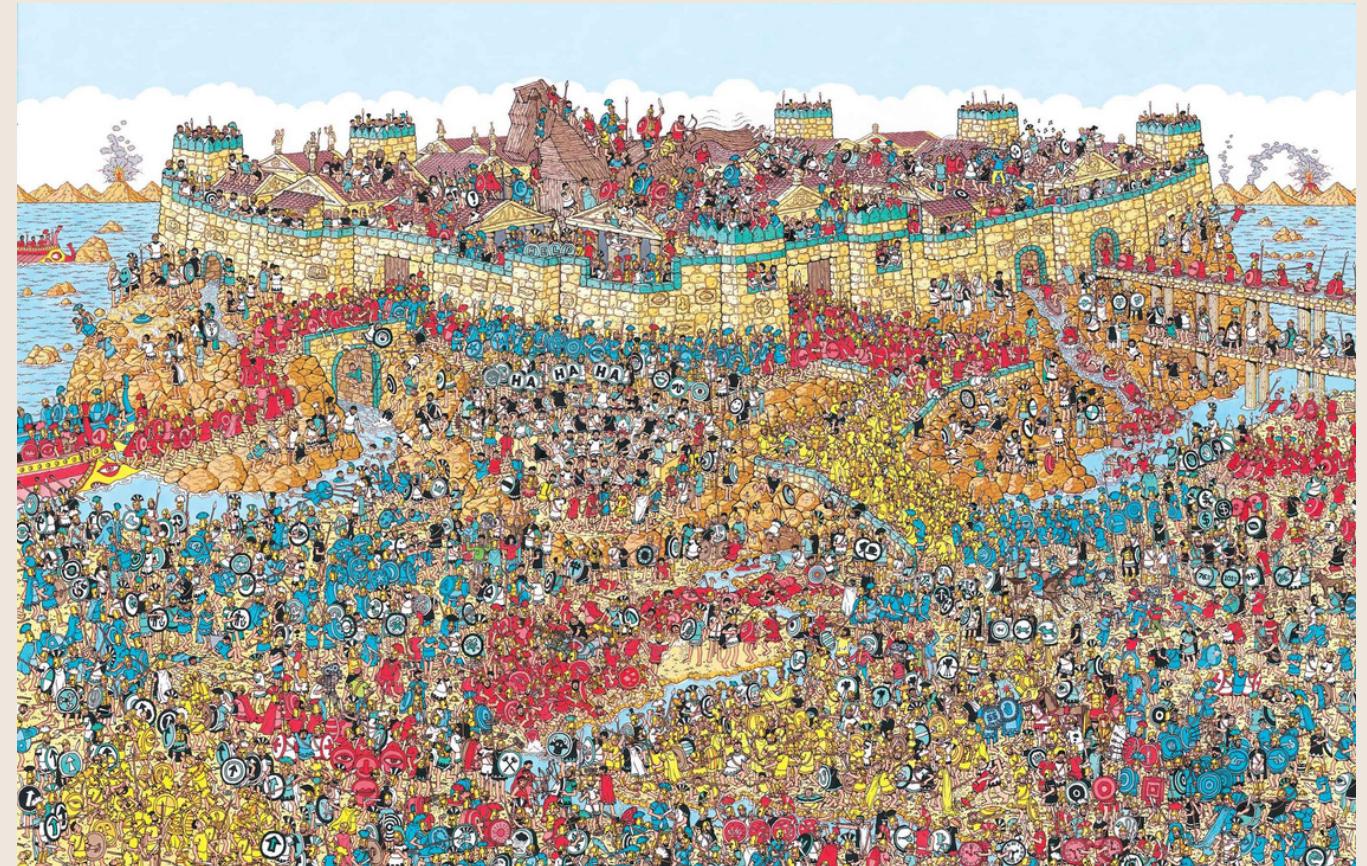
Not Waldo

~6000 images



Original Images

19 Images
containing Waldo



Uploading and Converting Images

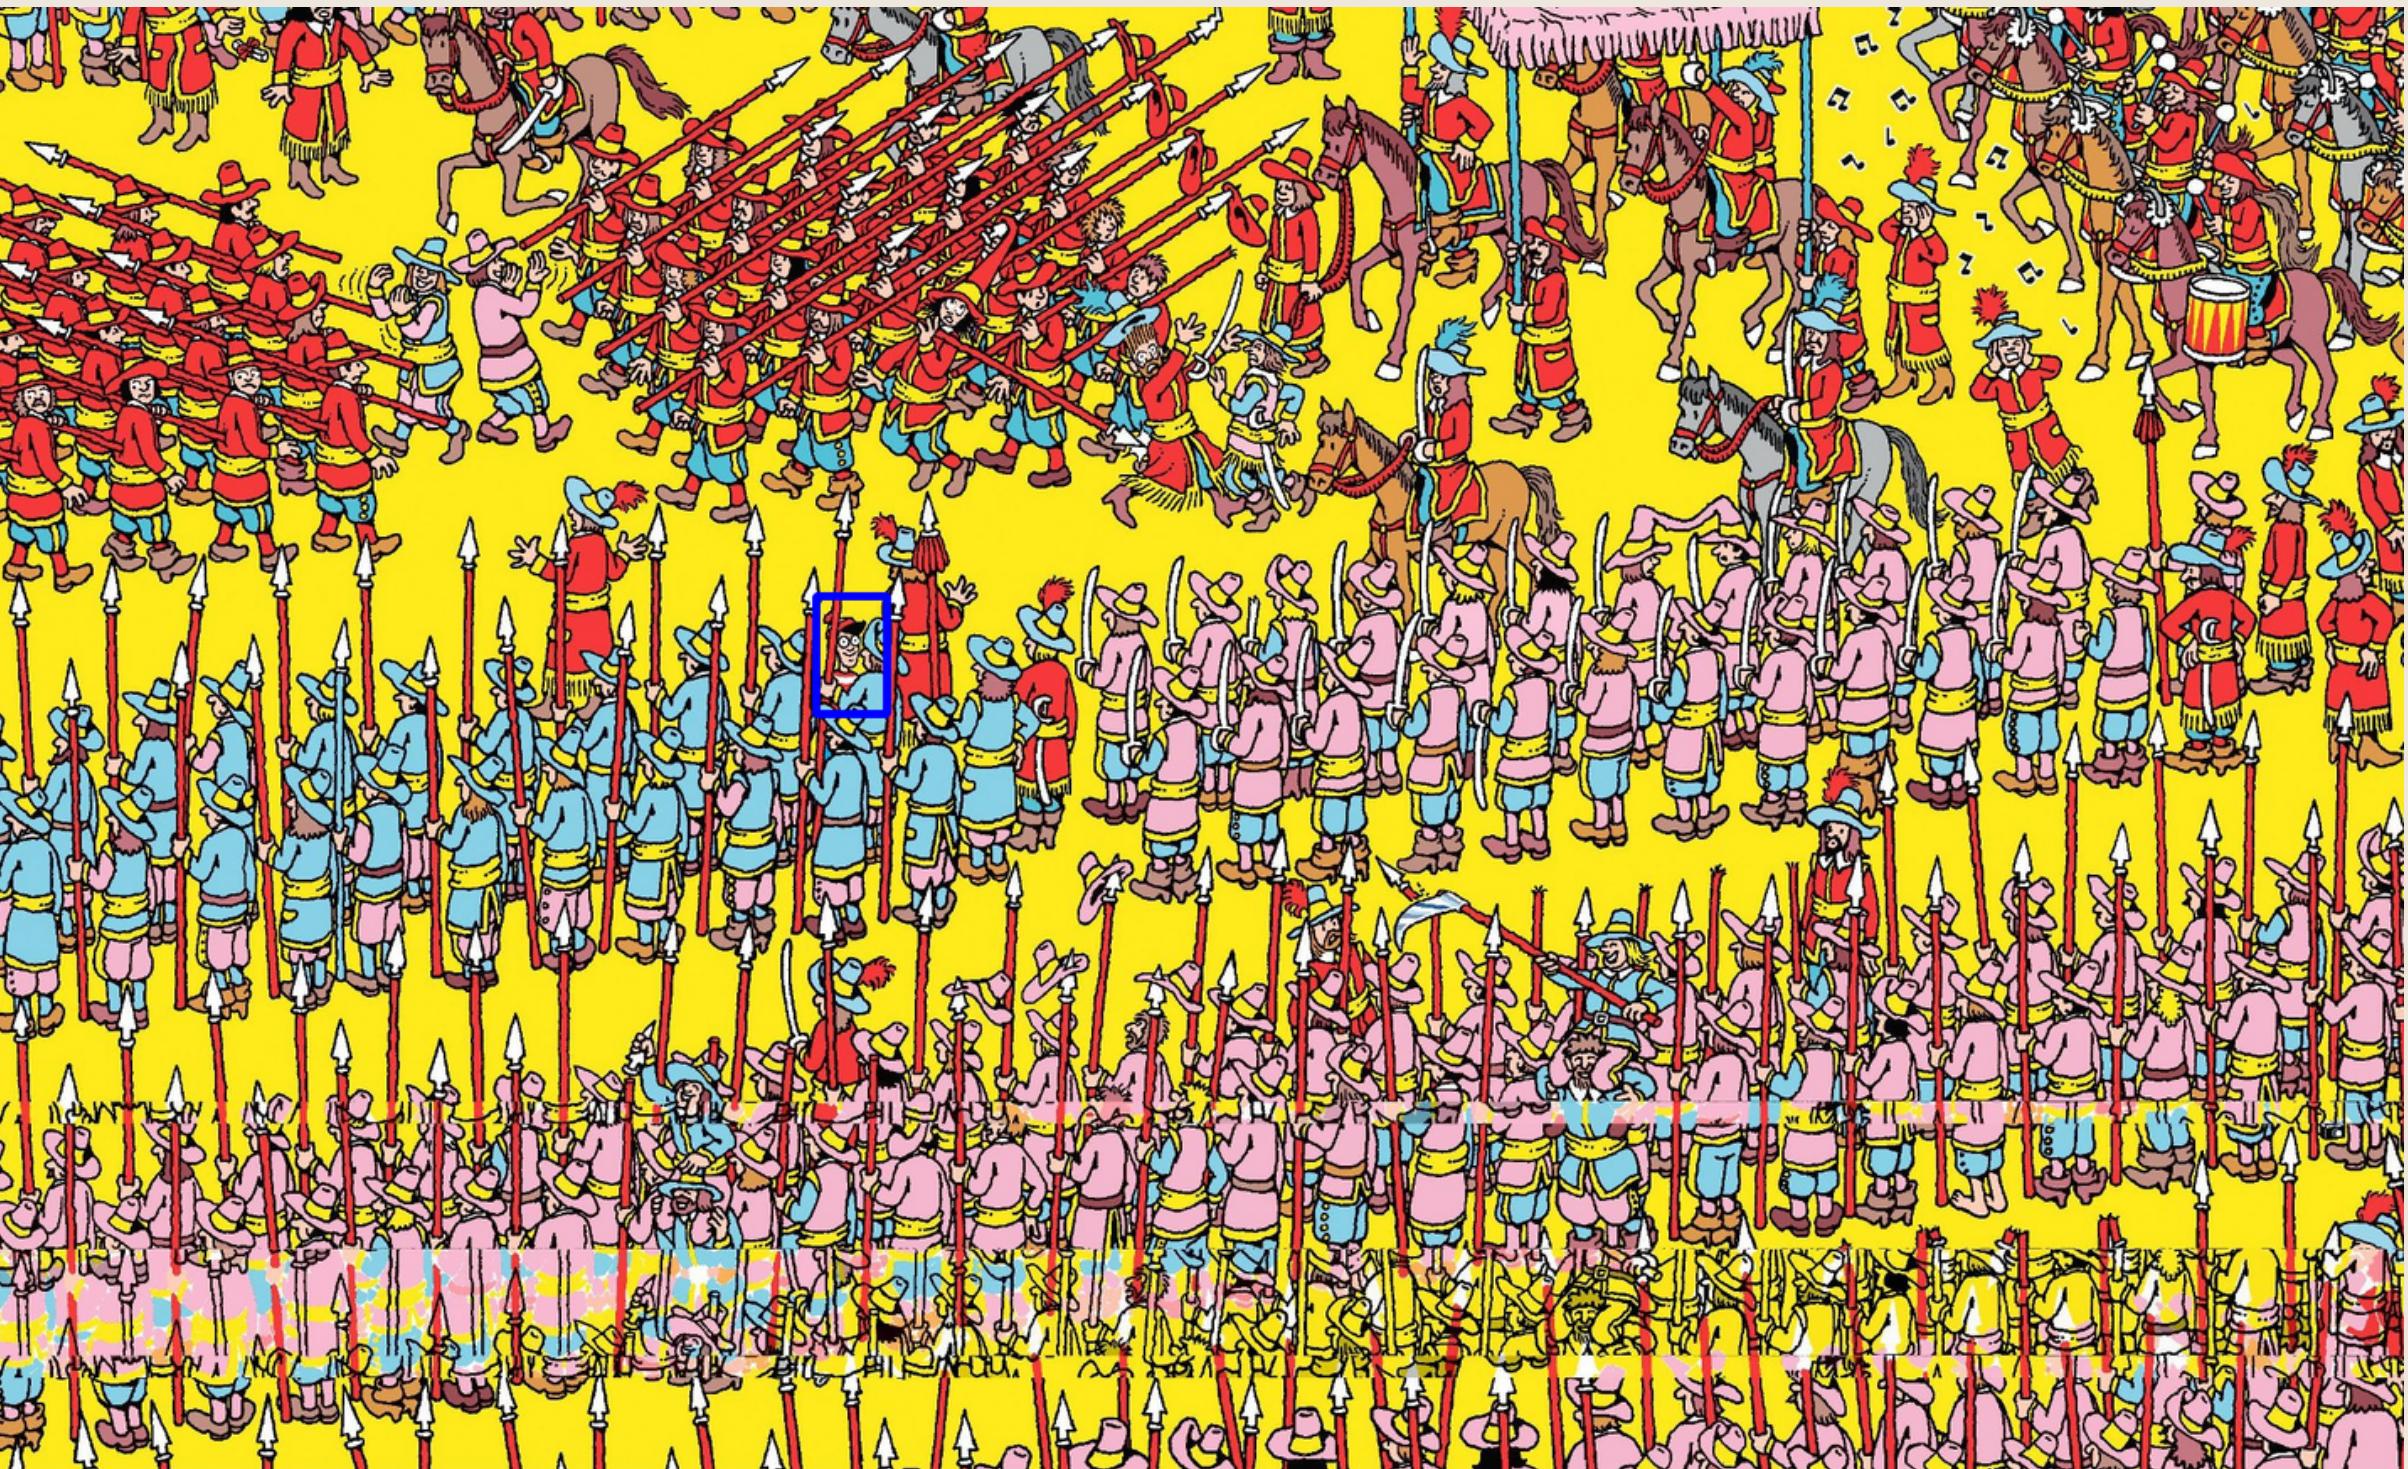
- cv2.imread to uploaded our images
- Converted into arrays of numbers (0-256)

```
[[[ 25  26  12]
 [ 25  26  12]
 [ 25  26  12]
 ...
 [113 106  52]
 [113 106  52]
 [114 107  53]]]
```



Annotating an Image

- Find bonding values on Waldo
- Annotating box
- Saving image
- Saving array of image into a csv file



Model





Image Classification

- Classifying images (Waldo =1, Not Waldo=0)
- Train and Test split
- Normalized Data
- Weighted Binary Crossentropy



Image Classification

Model Building

```
NUM_UNITS = 64
WINDOW_SHAPE = (3, 3)

model = Sequential()

# 2 x 64 convolutional neural network
model.add(Conv2D(NUM_UNITS, WINDOW_SHAPE, input_shape=X_train.shape[1:]))
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(NUM_UNITS, WINDOW_SHAPE))
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size=(2, 2)))

# flatten to 1D
model.add(Flatten())
model.add(Dense(64))

# output layer
model.add(Dense(1))
model.add(Activation("sigmoid"))

# compiling model
# create_weighted_binary_crossentropy(zero_weight, one_weight)
model.compile(loss="binary_crossentropy", optimizer="adam", metrics=["accuracy", tf.keras.metrics.Recall()], loss_weights=[zero_weight, one_weight]) # "binary_crossentropy"
```



Image Classification Model Summary

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_2 (Conv2D)	(None, 62, 62, 64)	1792
activation_3 (Activation)	(None, 62, 62, 64)	0
max_pooling2d_2 (MaxPooling 2D)	(None, 31, 31, 64)	0
conv2d_3 (Conv2D)	(None, 29, 29, 64)	36928
activation_4 (Activation)	(None, 29, 29, 64)	0
max_pooling2d_3 (MaxPooling 2D)	(None, 14, 14, 64)	0
flatten_1 (Flatten)	(None, 12544)	0
dense_2 (Dense)	(None, 64)	802880
dense_3 (Dense)	(None, 1)	65
activation_5 (Activation)	(None, 1)	0
<hr/>		
Total params: 841,665		
Trainable params: 841,665		
Non-trainable params: 0		

R-CNN

- Building a Box
- Identifying the Bounds
- Used to define the location of the target object
- Training Model on Bounds
- IoU
 - 1 Better
 - 0 Worse

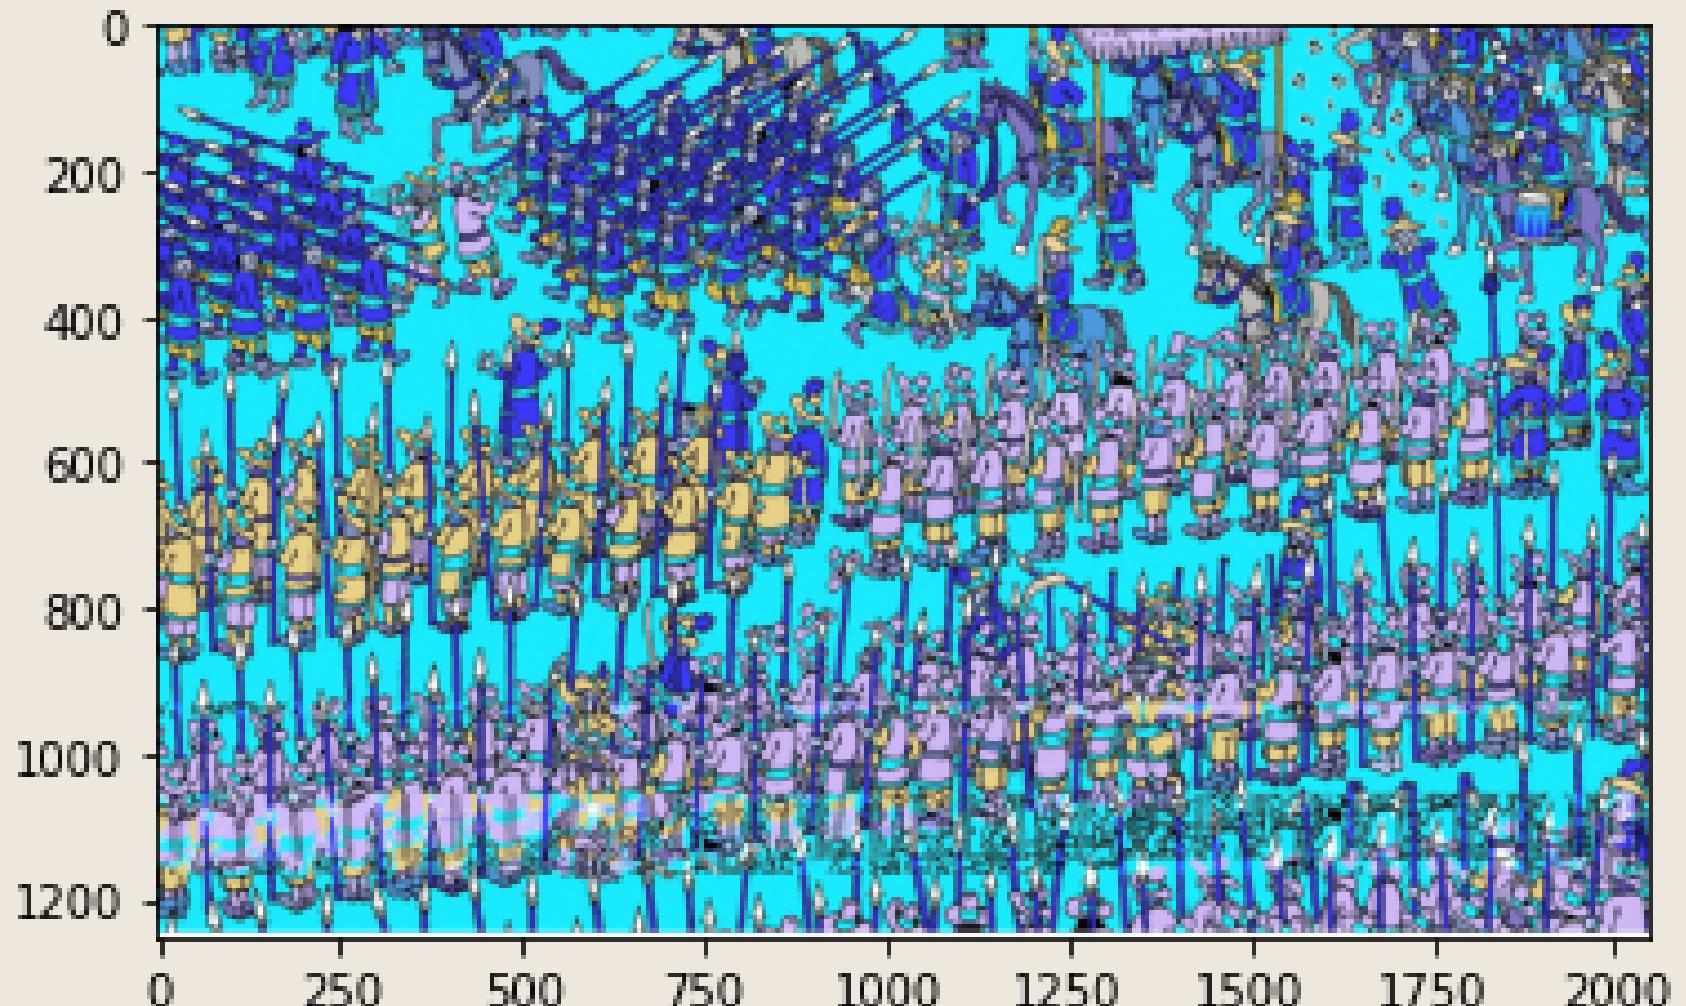




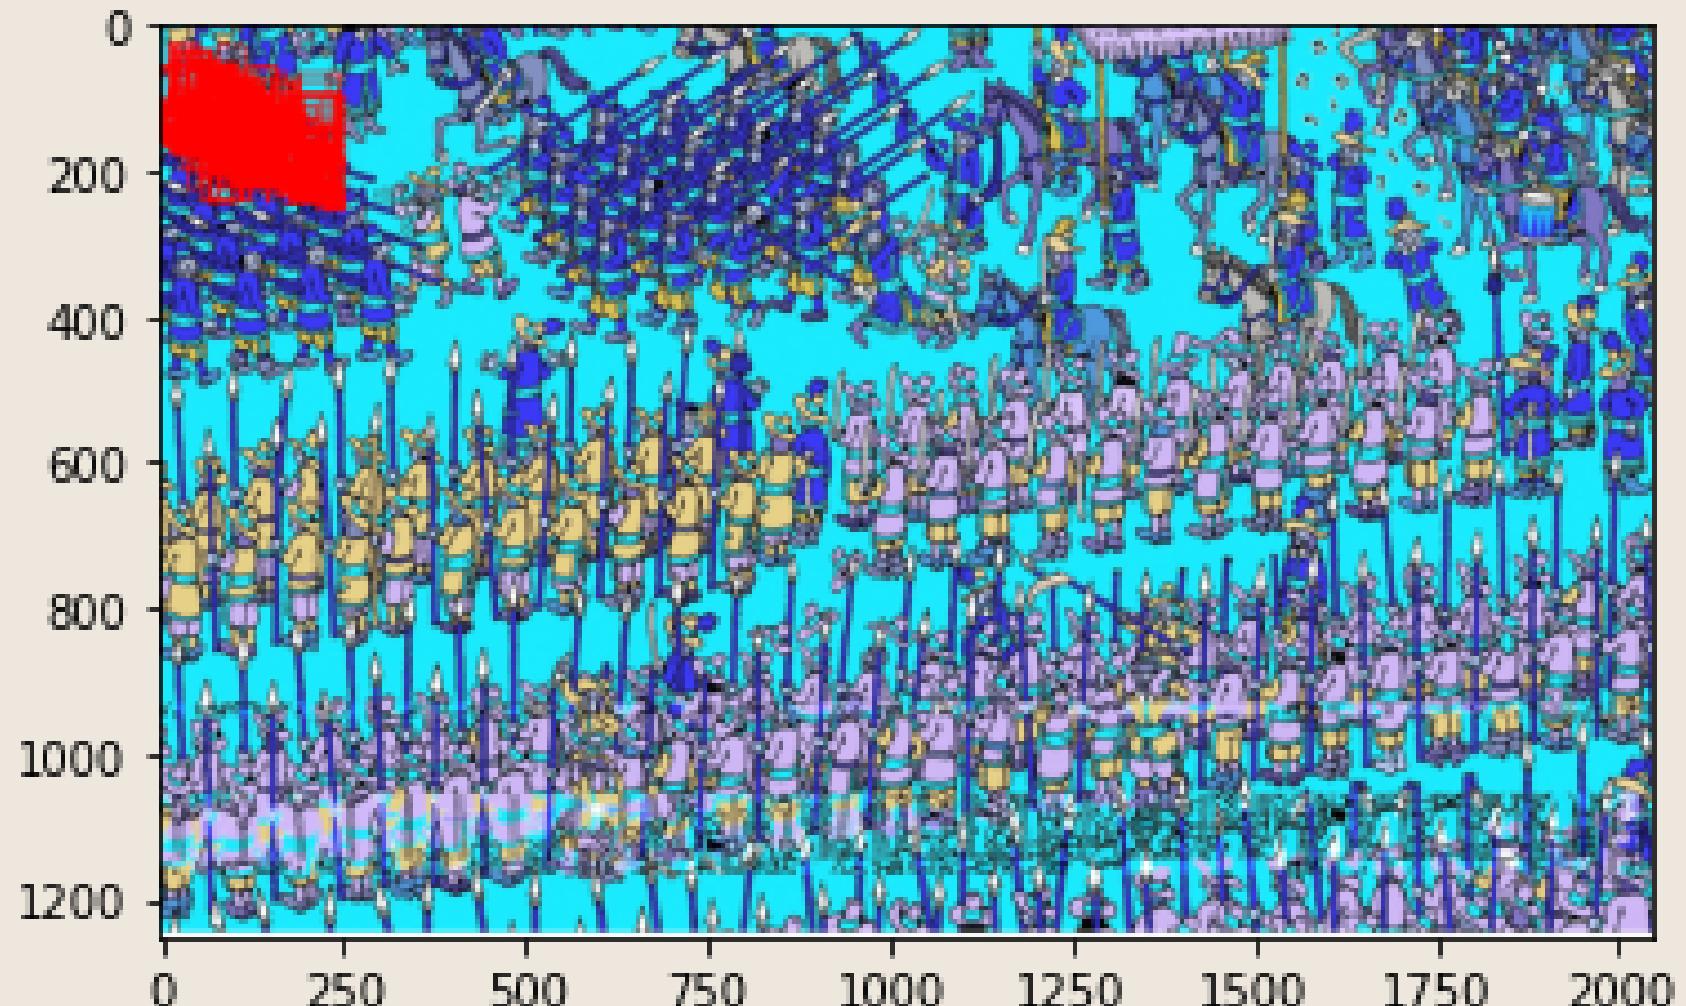
Labelled Image

```
for e,i in enumerate(os.listdir(original_img)):
    if e < 10:
        filename = i.split(".")[0]+".jpg"
        Index = 19
        filename2 = "waldo_"+str(Index)+".png.csv"
        img = cv2.imread(os.path.join(original_img,filename))
        #print(img)
        df = pd.read_csv(os.path.join(annot,filename2))
        plt.imshow(img)
        for row in df.iterrows():
            x1 = int(''.join(filter(str.isdigit, row[1][0].split(" ")[0])))
            y1 = int(''.join(filter(str.isdigit, row[1][0].split(" ")[1])))
            x2 = int(''.join(filter(str.isdigit, row[1][0].split(" ")[2])))
            y2 = int(''.join(filter(str.isdigit, row[1][2].split(" ")[1])))
            cv2.rectangle(img,(x1,y1),(x2,y2),(255,0,0), 2)
        plt.figure()
        plt.imshow(img)
        break
```

Original Image



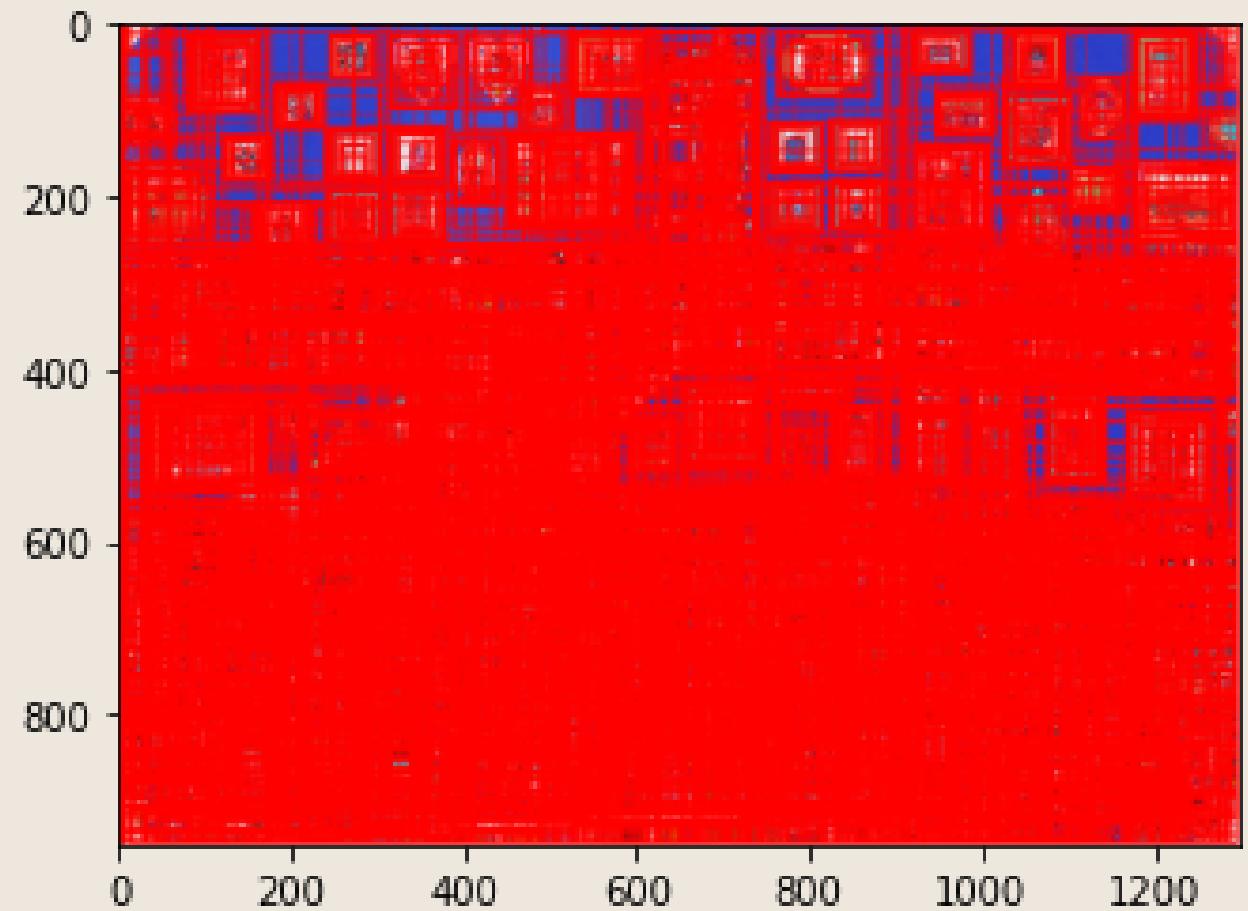
Annotated (not correct)



Selective Search

- Gets region proposals from selective search algorithm
- One single image
- Image on what the algorithm is doing

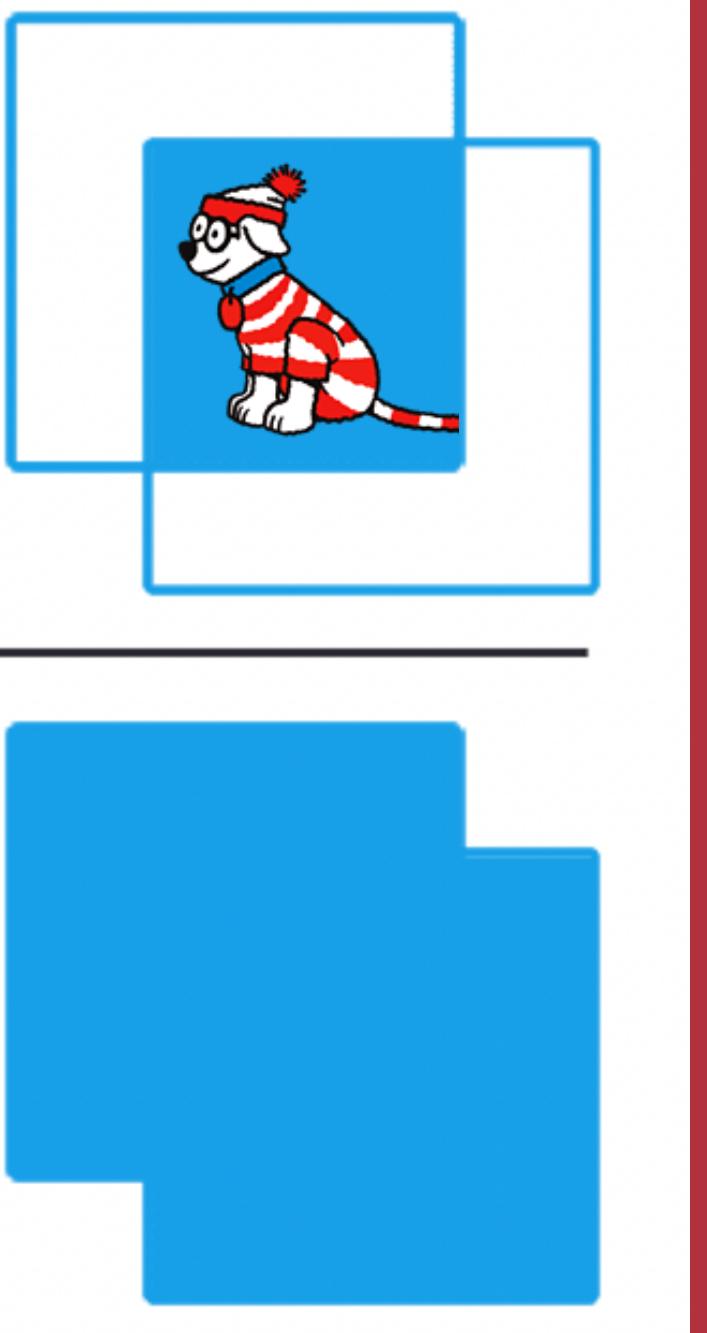
```
im = cv2.imread(os.path.join(original_img,'9.jpg'))
#cv2.imread(os.path.join(annot,"9.jpg"))
ss.setBaseImage(im)
ss.switchToSelectiveSearchFast()
rects = ss.process()
imOut = im.copy()
for i, rect in enumerate(rects):
    x, y, w, h = rect
    cv2.rectangle(imOut, (x, y), (x+w, y+h), (255, 0, 0), 1, cv2.LINE_AA)
    print(imOut)
plt.imshow(imOut)
```



Intersection Over Union

- Predicted bounding boxes from model.
- Boxes IoU greater than 0.7 (original paper it's 0.5) are considered positive example.
- Boxes with relative low IoU 0.3 negative examples.

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$



Conclusions



Classifying Waldo

- Our model can classify if there is or isn't Waldo even with unbalanced data
- Accuracy: ~80 %
- Results are inconsistent

Finding Waldo

- Annotation was a success.
- Since our index was out of bounds IoU could not identify the bounding boxes.
- Did not find Waldo.



Thank you!